



Contents lists available at ScienceDirect

Journal of Symbolic Computation

www.elsevier.com/locate/jsc

Faster arithmetic for number-theoretic transforms

David Harvey

University of New South Wales, Sydney, NSW, 2052, Australia



ARTICLE INFO

Article history:

Received 12 June 2013

Accepted 13 September 2013

Available online 24 September 2013

Keywords:

Number-theoretic transform

Fast Fourier transform

Modular arithmetic

Efficient algorithm

ABSTRACT

We show how to improve the efficiency of the computation of fast Fourier transforms over \mathbb{F}_p where p is a word-sized prime. Our main technique is optimisation of the basic arithmetic, in effect decreasing the total number of reductions modulo p , by making use of a redundant representation for integers modulo p . We give performance results showing a significant improvement over Shoup's NTL library.

© 2013 Elsevier B.V. All rights reserved.

1. Introduction

An important problem in computational number theory and cryptography is the efficient implementation of modular arithmetic. A typical implementation strategy is to represent elements of $\mathbb{Z}/N\mathbb{Z}$ by residues in a standard interval, such as $[0, N)$ or $[-N/2, N/2)$, and to reduce intermediate results to this interval after each operation in $\mathbb{Z}/N\mathbb{Z}$, such as addition or multiplication.

In many algorithms one can obtain a substantial constant factor speedup by delaying the modular reduction until after several arithmetic operations have been performed, by taking into account the bit-size of intermediate results. For example, to compute a dot product $\sum_i a_i b_i$, a fundamental operation in linear algebra, one may accumulate terms in batches, using ordinary integer (or even floating-point) arithmetic, and perform the reduction modulo N after each batch (Dumas et al., 2002).

The aim of this paper is to describe a strategy for reducing the number of modular reductions in the computation of a discrete Fourier transform over a finite field, also known as a *number-theoretic transform* (NTT). The NTT has a vast range of applications; we mention here only fast multiplication of large integers or polynomials (von zur Gathen and Gerhard, 2003).

For simplicity we restrict to the following situation. Let β describe the machine word size, for example $\beta = 2^{32}$ or $\beta = 2^{64}$. We work over $\mathbb{F}_p = \mathbb{Z}/p\mathbb{Z}$ where p is a word-sized prime, and we assume

that $p = 1 \pmod{L}$, where $L = 2^\ell$ is the transform length. This ensures that \mathbf{F}_p contains a primitive L -th root of unity; we fix one of these, denoted ω . The NTT is the map $(\mathbf{F}_p)^L \rightarrow (\mathbf{F}_p)^L$ defined by

$$b_j = \sum_{i=0}^{L-1} \omega^{ij} a_i, \quad 0 \leq j < L.$$

Equivalently, this is the map that evaluates the polynomial $a_0 + a_1 x + \cdots + a_{L-1} x^{L-1}$ at the points $1, \omega, \omega^2, \dots, \omega^{L-1}$.

It is well known that the fast Fourier transform (FFT) can be used to compute the NTT using $O(L \log L)$ operations in \mathbf{F}_p . For completeness, a simple in-place iterative radix-2 FFT algorithm is shown in [Algorithm 1](#).

Algorithm 1: Simple FFT

Input: $\omega \in \mathbf{F}_p$ and $x = (x_0, \dots, x_{L-1}) \in (\mathbf{F}_p)^L$, with $L = 2^\ell$
Output: DFT of x with respect to ω , in bit-reversed order

```

1 for i ← 1, 2, ..., ℓ do
2   |   ζ ← ω2i-1
3   |   m ← 2ℓ-i
4   |   for 0 ≤ j < 2i-1 do
5   |   |   t ← 2jm
6   |   |   for 0 ≤ k < m do
7   |   |   |   [xt+k | xt+k+m] ← [xt+k + xt+k+m | ζk(xt+k - xt+k+m)]
8   |   |   end
9   |   end
10 end

```

Our main focus in this paper is on the *butterfly operation* in line 7, which computes $(X, Y) \mapsto (X + Y, W(X - Y))$ for some fixed root of unity $W \in \mathbf{F}_p$ (a suitable power of ω). At first glance this requires one modular addition, one modular subtraction, and one modular multiplication, and this is how the butterfly is usually implemented. Of course these modular operations are themselves implemented on modern microprocessors using more basic primitives. For example, a modular addition is usually implemented as an ordinary integer addition, followed by a comparison with p , followed by a conditional subtraction. In this paper we investigate the butterfly at this lower level, showing how to streamline the implementation to reduce the number of comparisons and conditional operations. Another interpretation is that we have reduced the number of modular reductions.

2. A typical butterfly implementation

Victor Shoup's NTL (Number Theory Library) ([Shoup, 2013](#)) is a popular C++ library used in computational number theory. It makes heavy use of the NTT. Its implementation of the butterfly $(X, Y) \mapsto (X + Y, W(X - Y))$ may be expressed by the pseudocode shown in [Algorithm 2](#). It has been simplified to focus attention on the arithmetic aspects, ignoring issues like loop unrolling, software pipelining, and locality. All variables represent register-sized quantities.

Note that W and W' are independent of the data being transformed; they can be precomputed and reused for each transform. A similar comment applies to all the butterfly algorithms considered in this paper.

Theorem 1. [Algorithm 2](#) is correct.

Proof. Lines 1–2 compute the sum $X + Y$ and reduce it modulo p to the standard interval $[0, p)$, using the assumption $p < \beta/2$ to avoid overflow in the first line. Lines 3–4 compute $T = X - Y \pmod{p}$ in the same way, assuming that T has a signed type for the comparison.

Algorithm 2: NTL butterfly implementation

Input: $p < \beta/2$
 $0 < W < p$
 $W' = \lfloor W\beta/p \rfloor$, $0 < W' < \beta$
 $0 \leq X < p$
 $0 \leq Y < p$

Output: $X' = X + Y \pmod{p}$
 $Y' = W(X - Y) \pmod{p}$

- 1 $X' \leftarrow X + Y$
- 2 **if** $X' \geq p$ **then** $X' \leftarrow X' - p$
- 3 $T \leftarrow X - Y$
- 4 **if** $T < 0$ **then** $T \leftarrow T + p$
- 5 $Q \leftarrow \lfloor W'T/\beta \rfloor$
- 6 $Y' \leftarrow (WT - Qp) \pmod{\beta}$
- 7 **if** $Y' \geq p$ **then** $Y' \leftarrow Y' - p$
- 8 **return** X', Y'

Lines 5–7 compute the product $WT \pmod{p}$. Line 5 first generates an estimated quotient Q . By the definition of W' and Q we have

$$0 \leq \frac{W\beta}{p} - W' < 1, \quad 0 \leq \frac{W'T}{\beta} - Q < 1.$$

Multiplying by Tp/β and p respectively, and adding, yields

$$0 \leq WT - Qp < \frac{Tp}{\beta} + p < 2p < \beta.$$

In particular, line 6 correctly computes $Y' = WT - Qp$, and the single correction in line 7 suffices to reduce it into $[0, p)$. \square

The operations performed by Algorithm 2 map comfortably onto modern instruction sets. In line 5, the expression $\lfloor W'T/\beta \rfloor$ represents the high word of the product $W'T$, which can typically be obtained by a single hardware multiply instruction. Both WT and Qp in line 6 require only the low word of the product. The conditional additions and subtractions in lines 2, 4 and 7 are simple enough that a modern compiler will implement them by a conditional move instruction rather than by a branch.

If we assume that X and Y are distributed uniformly in $[0, p)$, then the first two conditions (lines 2 and 4) hold with probability 50%. The behaviour of the third condition (line 7) is more complex. If the quantity $W\beta/p - W'$ considered in the proof of the theorem happens to be close to zero, then the condition will be satisfied with very low probability. At the other extreme, if $W\beta/p - W'$ is close to unity, and if we assume that p is near $\beta/2$, then one can show the condition is satisfied with probability about 25%. In this case it is still reasonable to prefer a conditional move instruction to a branch.

The modular multiplication algorithm in lines 5–7 is attributed to Shoup in Färnqvist (2005), but does not seem to have been published. It first appears in NTL version 5.4 in 2005. The use of a suitable precomputed approximation to W/p implies that only a single correction step (line 7) is necessary, and that the remainder is obtained using only multiplication modulo β (line 6), an advantage on processors that can compute the low word faster than the full product. The latter idea is also used heavily in the division algorithm of Möller and Granlund (2011).

In our exposition we assumed for simplicity that $p < \beta/2$, but Niels Möller has pointed out (personal communication) that this can be improved. The modular subtraction can be made to work for any $p < \beta$ by replacing the condition $T < 0$ by $X < Y$, or indeed by using the borrow generated by the subtraction $T = X - Y$. The addition can be treated similarly by rewriting it as $X' = X - (p - Y)$. A more careful analysis of the candidate remainder $WT - Qp$ then shows that the entire algorithm works for any $p < \beta/\phi$ where $\phi = \frac{1}{2}(1 + \sqrt{5}) \approx 1.618$.

3. A modified butterfly

In this section we propose several modifications to [Algorithm 2](#). Our motivation is that the adjustment steps in lines 2, 4 and 7 are relatively expensive on modern microprocessors, compared to hardware integer multipliers, which in recent years have become very fast.

Our first observation is that Shoup's algorithm for computing $WT \pmod{p}$ does not require that $0 \leq T < p$; in fact the proof given above shows that it works for any $0 \leq T < \beta$. Therefore we may replace lines 3–4 by simply $T \leftarrow X - Y + p$, after which we have $0 \leq T < 2p < \beta$, and the algorithm proceeds as before. This simplification is not new, although it does not appear to be well known. It is not used in NTL. It was apparently used by Fabrice Bellard in a computation of π to 2.7 trillion decimal places in 2009 (personal communication).

Our second observation appears to be new, and is the main novelty introduced in the present paper. Namely, we may also remove the adjustment in line 7, provided that throughout the FFT we use a *redundant* representation for elements of \mathbb{F}_p . That is, instead of representing elements of \mathbb{F}_p by integers in $[0, p)$, we use the wider interval $[0, 2p)$, so each element has two possible representations. For this to work, we must modify the butterfly to accept *inputs* in $[0, 2p)$, and we must impose the stronger condition $p < \beta/4$. Pseudocode for the resulting butterfly is shown in [Algorithm 3](#).

According to Jason Papadopolous (personal communication), around 1998 Ernst Mayer suggested the use of a redundant representation in the context of a fast Galois transform, i.e. a Fourier transform over \mathbb{F}_{q^2} where $q = 2^{61} - 1$. Our new algorithm may be regarded as a generalisation of this idea to the case of an NTT with arbitrary modulus p .

Algorithm 3: Modified Shoup butterfly

Input: $p < \beta/4$
 $0 < W < p$
 $W' = \lfloor W\beta/p \rfloor$, $0 < W' < \beta$
 $0 \leq X < 2p$
 $0 \leq Y < 2p$

Output: $X' = X + Y \pmod{p}$, $0 \leq X' < 2p$
 $Y' = W(X - Y) \pmod{p}$, $0 \leq Y' < 2p$

- 1 $X' \leftarrow X + Y$
- 2 **if** $X' \geq 2p$ **then** $X' \leftarrow X' - 2p$
- 3 $T \leftarrow X - Y + 2p$
- 4 $Q \leftarrow \lfloor W'T/\beta \rfloor$
- 5 $Y' \leftarrow (WT - Qp) \pmod{\beta}$
- 6 **return** X', Y'

Theorem 2. [Algorithm 3](#) is correct.

Proof. Lines 1–2 compute a representative for $X + Y \pmod{p}$ in the interval $[0, 2p)$; this does not overflow as $p < \beta/4$. Line 3 computes a representative T for $X - Y \pmod{p}$ in the interval $[0, 4p)$. Lines 4–5 compute a representative for $WT \pmod{p}$ in the interval $[0, 2p)$, in the same way as in the proof of [Theorem 1](#). \square

Both outputs X', Y' lie in $[0, 2p)$, ready to be processed by a subsequent butterfly. Depending on the needs of the NTT user, it may be necessary to normalise the final NTT output into the interval $[0, p)$, imposing an additional $O(L)$ cost.

4. Variants

In some situations one requires a butterfly of the form $(X, Y) \mapsto (X + WY, X - WY)$. This appears if one switches from a 'decimation-in-frequency' transform to a 'decimation-in-time' transform. It also

appears naturally in the inverse FFT obtained by running [Algorithm 1](#) backwards. [Algorithm 4](#) shows an analogue of [Algorithm 3](#) for this case. The main difference is that elements of \mathbf{F}_p are represented by an integer in $[0, 4p)$, so each residue has four possible representatives. As in [Algorithm 3](#), this strategy saves two modular reductions compared to the usual implementation.

Algorithm 4: Modified inverse butterfly

Input: $p < \beta/4$
 $0 < W < p$
 $W' = \lfloor W\beta/p \rfloor, 0 < W' < \beta$
 $0 \leq X < 4p$
 $0 \leq Y < 4p$

Output: $X' = X + WY \pmod{p}, 0 \leq X' < 4p$
 $Y' = X - WY \pmod{p}, 0 \leq Y' < 4p$

- 1 **if** $X \geq 2p$ **then** $X \leftarrow X - 2p$
- 2 $Q \leftarrow \lfloor W'Y/\beta \rfloor$
- 3 $T \leftarrow (WY - Qp) \pmod{\beta}$
- 4 $X' \leftarrow X + T$
- 5 $Y' \leftarrow X - T + 2p$
- 6 **return** X', Y'

Theorem 3. [Algorithm 4](#) is correct.

Proof. Line 1 reduces X into $[0, 2p)$. Lines 2–3 compute $T = WY \pmod{p}$ in $[0, 2p)$, in the same way as [Algorithm 3](#). The remaining lines complete the calculation of X' and Y' in $[0, 4p)$. \square

The same idea may be applied to butterflies using other modular multiplication algorithms. [Algorithm 5](#) shows a variant using Montgomery multiplication ([Montgomery, 1985](#)). The idea of Montgomery multiplication is in effect to replace the usual Euclidean quotient by a 2-adic quotient. Compared to [Algorithm 3](#), this butterfly algorithm has the advantage that only a single value W' must be stored in a table for each root of unity (W is not actually used in the algorithm). On the other hand, one of the multiplications modulo β has been replaced by a full product, which may be more expensive on some processors.

Algorithm 5: Modified Montgomery butterfly

Input: p odd, $p < \beta/4$
 $0 < W < p$
 $W' = \beta W \pmod{p}, 0 < W' < p$
 $J = p^{-1} \pmod{\beta}$
 $0 \leq X < 2p$
 $0 \leq Y < 2p$

Output: $X' = X + Y \pmod{p}, 0 \leq X' < 2p$
 $Y' = W(X - Y) \pmod{p}, 0 \leq Y' < 2p$

- 1 $X' \leftarrow X + Y$
- 2 **if** $X' \geq 2p$ **then** $X' \leftarrow X' - 2p$
- 3 $T \leftarrow X - Y + 2p$
- 4 $R_1\beta + R_0 \leftarrow W'T$
- 5 $Q \leftarrow R_0J \pmod{\beta}$
- 6 $H \leftarrow \lfloor Qp/\beta \rfloor$
- 7 $Y' \leftarrow R_1 - H + p$
- 8 **return** X', Y'

Theorem 4. [Algorithm 5](#) is correct.

Table 1

Cycles per butterfly for several implementations.

	Westmere	Sandy Bridge	Piledriver
NTL	12.0	9.1	16.7
Algorithm 2	10.1	9.1	13.9
Algorithm 3	6.9	5.9	12.0

Proof. Lines 1–3 are identical to the corresponding lines in [Algorithm 3](#). Line 4 computes the product $W'T$, placing the low and high words of the result in R_0 and R_1 respectively, so that $0 \leq R_1 < p$. Line 5 computes $Q = R_0/p \pmod{\beta}$ (this may be regarded as a 2-adic approximation to the quotient $W'T/p$). Line 6 computes the high word of Qp , so that $0 \leq H < p$. We have $Qp = R_0 \pmod{\beta}$, so $Qp = H\beta + R_0$. Then $W'T - Qp = \beta(R_1 - H)$, and thus $WT = R_1 - H \pmod{p}$. This agrees modulo p with the value computed for Y in line 7, which lies in the interval $[0, 2p]$.

(Usually in Montgomery's algorithm, a further comparison and conditional subtraction/addition is performed to normalise the result into $[0, p)$, but we have simply skipped that.) \square

5. Implementation and performance

The practical benefit (if any) derived from the new algorithms depends heavily on the hardware used. To illustrate what may occur in practice, we performed some timing experiments on several machines: a 3.06 GHz Intel Xeon (model X5675, ‘Westmere’ microarchitecture), a 2.6 GHz Intel Xeon (model E5-2670, ‘Sandy Bridge’ microarchitecture), and an AMD Vishera (model FX-8350, ‘Piledriver’ microarchitecture), all modern 64-bit processors.

We compared a C implementation of a number-theoretic transform incorporating [Algorithm 3](#), a similar implementation of a transform based on [Algorithm 2](#), and the FFT routine in NTL itself (version 5.5.2). The code is available from the author's web site. It is moderately optimised: the inner loops are 2-way unrolled, and the last two layers of the FFT have dedicated loops. The values of W and W' are precomputed and stored in a table. Everything was compiled using GCC 4.6.3 on the Intel machines, and GCC 4.7.2 on the AMD machine, with the $-O2$ optimisation flag. We also tried the $-O3$ flag; this made no significant difference to the results. Access to the high word of the product of two 64-bit integers is obtained using the compiler's built-in support for 128-bit integer types; in NTL this is handled via inline assembly macros.

We ran transforms of length $2^{11} = 2048$. This is long enough to avoid too much loop overhead, but short enough that all memory accesses hit the L1 cache, so we may ignore locality problems. For our code we used a 62-bit prime p , and for NTL we used a 50-bit prime. NTL supports moduli of only up to 50 bits on a 64-bit machine, for historical reasons related to floating-point arithmetic; however, only integer arithmetic is used in the actual FFT routine.

The NTL FFT routine also performs a bit-reversal of the input array and computes a table of roots of unity on each FFT invocation. To make a fair comparison with our code, which does not perform these steps, we removed their contribution in the timing data shown in [Table 1](#).

The table shows the total FFT time (CPU clock cycles) divided by the number of butterflies performed, which is 11×2^{10} . Note that $O(L)$ of the $O(L \log L)$ butterflies have $W = 1$, and these are faster than the general butterfly with $W \neq 1$. Therefore these figures slightly underestimate the cost of each general butterfly.

Two features of the table are worth pointing out. First, our implementation of [Algorithm 2](#) is competitive with NTL. In fact, on two of the three platforms, it is slightly faster. Second, [Algorithm 3](#) outperforms [Algorithm 2](#) by a factor of about 1.5 on the Intel machines and about 1.15 on the AMD machine.

The Westmere processor can sustain a maximum throughput of one 64-bit multiplication every 2 cycles ([Fog, 2013](#)). If we assume that a butterfly requires three such multiplications, then the best we can expect is 6 cycles per butterfly. The reported figure of 6.9 cycles comes fairly close to this; in other words, the multiplier is close to saturated. For Sandy Bridge, the relevant figure is either 1 cycle or 2 cycles per multiplication, depending on whether one of the operands is in a register or

fetched from L1 cache. In this case the limiting factor is more likely to be cache bandwidth. Finally, for Piledriver the throughput is 4 cycles per multiplication, so again we are close to saturating the multiplier.

Acknowledgements

The author thanks Tommy Färnqvist, Torbjörn Granlund, Niels Möller, Jason Papadopolous and Paul Zimmermann for stimulating conversations on this topic, and Torbjörn Granlund for providing access to the AMD machine. The referees provided many helpful comments. The author was partially supported by the Australian Research Council, DECRA Grant DE120101293.

References

- Dumas, Jean-Guillaume, Gautier, Thierry, Pernet, Clément, 2002. Finite field linear algebra subroutines. In: Proceedings of the 2002 International Symposium on Symbolic and Algebraic Computation. ACM, New York, pp. 63–74 (electronic). MR 2035233.
- Färnqvist, Tommy, 2005. Number theory meets cache locality – efficient implementation of a small prime FFT for the GNU Multiple Precision Arithmetic Library. Master's thesis. Stockholm University. http://www.nada.kth.se/utbildning/grukth/exjobb/rapportlistor/2005/rapporter05/farnqvist_tommy_05091.pdf.
- Fog, Agner, 2013. Instruction tables: Lists of instruction latencies, throughputs and micro-operation breakdowns for Intel, AMD and VIA CPUs. <http://www.agner.org/optimize/>.
- Möller, Niels, Granlund, Torbjörn, 2011. Improved division by invariant integers. *IEEE Trans. Comput.* 60, 165–175.
- Montgomery, Peter L., 1985. Modular multiplication without trial division. *Math. Comput.* 44 (170), 519–521. MR 777282 (86e:11121).
- Shoup, Victor, 2013. NTL: a library for doing number theory (Version 5.5.2). <http://www.shoup.net/ntl/>.
- von zur Gathen, Joachim, Gerhard, Jürgen, 2003. Modern Computer Algebra, second ed. Cambridge University Press, Cambridge. MR 2001757 (2004g:68202).